


- What progress have you made so far? (share any preliminary results)

I have completed the following tasks:

1. Data collection
2. Data Preparation: The original dataset contained nearly 450 images and corresponding annotation files. The annotation files were in the form of .xml files which contained information about the images and multiple bounding boxes for each image. I used the xml library's etree module to parse the .xml files and extract the bounding box coordinates and the class labels. I then analyzed the various features and removed several features that were constant accross all images since these were unlikely to have much predictive power. I also detected several bounding boxes that were outside the image bounds and removed these. This resulted in the removal of about 20 bounding boxes from the expected 1450. I stored the information about the overall image in a file called `images.csv` and that about the bounding boxes in a file called `objects.csv`. `objects.csv` contains a foreign key mapping all bounding boxes to the images they belong to in `images.csv`.
3. Image Snippet Extraction: Since we need to label each bounding box, I extracted the part of the image that was contained within the bounding box and saved it as a separate image. The image `BikesHelmets0.png` is shown below. Followed by the snippets extracted from it.



Image BikesHelmets0.png



4. Snippet Resizing: Since I plan on using a CNN for this task, I had to resize all the images of the areas of interest to the same size. I used some visual data analysis

using scatterplots of image resolutions to get a fair estimate of the ideal size. The original scatterplot was dense in certain regions, so it was hard to get a good estimate and had to be zoomed into. I determined that a size of about 45x45 pixels would be a decent place to start.

- Some of the challenges I have so far have primarily been due to the differences between the `torchvision` library and `sklearn` which is what I have predominantly used in the past. However, the `torchvision` library is very well documented and I am beginning to grasp how to use transformers and define custom ones for various stages of the pipeline.
- Here is a tentative timeline for the rest of the project:
 1. **December 3rd:** Complete implementing the transformer steps for image resizing, initial model training and testing.
 2. **December 4th-5th:** Experiment with different model architectures and hyperparameters to improve performance.
 3. **December 6th-7th:** Finalize the model and generate the final outputs.
 4. **December 8th:** Prepare the final report and submit the project.

The code for the project has been sent to the email address provided in the instructions.